


Dutch Sentiment Classification — Experiment Summary

LLM leads overall; Jina Ordinal leads research; TF-IDF remains deployable

4,800 raw → audit + dedupe → 3,838 train samples + 960 test samples



OPERATING METRICS + MODEL GOVERNANCE

MODEL	ACCURACY	NEGATIVE P / R (n=60)	ROLE	DECISION / CONSTRAINT
LLM few-shot (24-shot)	72.08%	77.46% / 91.67%	External advisor	Quality leader • external API
Jina Ordinal	68.96%	72.73% / 80.00%	Research only	Best research score • non-commercial
Jina Logistic	67.19%	54.08% / 88.33%	Research only	Frozen encoder + trained head
RobBERT Improved Ensemble	67.71%	69.39% / 56.67%	Challenger	Best RobBERT • Drive weights
TF-IDF Ordinal	65.00%	63.79% / 61.67%	Challenger	Ordinal local alternative
TF-IDF Logistic	65.31%	72.09% / 51.67%	Production benchmark	Deployed CPU champion
RobBERT v2 Logistic	63.54%	58.73% / 61.67%	Test evidence	Initial end-to-end baseline

BEST OVERALL: LLM FEW-SHOT • DEPLOYABLE: TF-IDF • BEST END-TO-END TRANSFORMER: ROBBERT IMPROVED

Lower than ~90% literature benchmarks: smaller, imbalanced 3-class data with an ambiguous Average label—scored by Macro-F1, not balanced binary accuracy.